

Temperature Fragility and the Conditional Benefits of Truncation Sampling

Anonymous ACL submission

Abstract

Large language models generate text by sampling each token from a predicted distribution, and a temperature parameter sets how far the draw strays from the most probable tokens. Truncation samplers such as top- p and min- p discard the least probable tokens before the draw, so that sampling at high temperature stays coherent. Their reported accuracy gains come from temperatures of 1.5 to 3, while the defaults of deployed systems cluster between 0.6 and 1.0. Whether they change accuracy at those defaults, and for which models, has not been measured. We test thirteen open-weight models on GSM8K and MMLU-Pro at temperatures 0.7, 1.0, and 1.3 in one controlled pipeline, seven of them under eight decoding configurations. Six of the thirteen models lose 17 to 38 accuracy points on MMLU-Pro between 0.7 and 1.3, and the other seven lose at most 10. The lost accuracy comes from generations that run to the token limit or never state an answer. These results suggest that truncation samplers improve accuracy primarily when higher temperatures substantially degrade model performance. Where accuracy remains stable across temperatures, none of the tested truncation samplers improves on plain temperature sampling.

1 Introduction

Sampling is the default way deployed language models generate text (Shi et al., 2024). At each step the model assigns a probability to every token in its vocabulary, and the decoder draws one from that distribution. A temperature parameter reshapes the distribution before the draw. Low temperatures concentrate probability on the few most likely tokens. High temperatures spread it toward the tail of the distribution, the many tokens the model rates as unlikely, where its estimates are least reliable (Holtzman et al., 2020; Hewitt et al., 2022). The typical defaults of model vendors and inference engines cluster between 0.6 and 1.0 (Renze and

Guven, 2024; Hochlehnert et al., 2025),¹ and every engine implements truncation samplers beside them. Truncation samplers, among them top- p , min- p , and top- $n\sigma$, discard the tail of the distribution before the draw, each by its own criterion. They were designed to keep open-ended text coherent at high temperature (Garces Arias et al., 2025), and they are now also evaluated on tasks with one checkable answer, where the claim is accuracy: min- p reports gains on GSM8K (Nguyen et al., 2025), and top- $n\sigma$ reports accuracy held up to a temperature of 3.0 (Tang et al., 2024).

Both results rest on one model and one range of temperatures. Top- $n\sigma$ was evaluated on Llama-3-8B (Tang et al., 2024) and min- p on Mistral-7B (Nguyen et al., 2025), at temperatures of 1.5 to 3, above the defaults systems set. A reanalysis that searched the same number of hyperparameter settings for every rule found the gains of min- p indistinguishable from those of the other rules (Schaeffer et al., 2025). The question cannot be settled by comparing across papers, because benchmark scores move with the inference backend even under greedy decoding (Masoudian et al., 2026), and with the prompt, the parser, and the item set. Settling it requires one pipeline in which everything except the decoding configuration is held fixed. It also requires enough models to show whether the answer depends on the model.

Research question: Under what model and temperature conditions do truncation samplers improve accuracy over plain temperature sampling?

We test thirteen open-weight models from the Llama-3, Mistral, Qwen, Gemma, and OLMo fam-

¹generation_config.json of meta-llama/Llama-3.1-8B-Instruct sets 0.6; llama.cpp’s common/sampling.h sets 0.8; the OpenAI chat-completions and Anthropic messages API references set 1.0. All read July 2026. Recommendations outside this range exist: Mistral’s instruct model cards suggest 0.15.

ilies on GSM8K and MMLU-Pro at temperatures 0.7, 1.0, and 1.3, in one pipeline with fixed items, prompts, token budgets, and parser. Ten of the models also run under eight decoding configurations: greedy decoding, plain temperature sampling, and six truncation configurations. Section 3 gives the design. Two questions follow. The first is whether a model keeps its accuracy between 0.7 and 1.3 under plain sampling, and what the lost accuracy looks like when it does not. The second is whether any truncation rule improves on plain temperature sampling between 0.7 and 1.3.

We find that temperature sensitivity differs by model. Between 0.7 and 1.3, six of the thirteen models lose 17 to 38 accuracy points on MMLU-Pro under plain sampling, and the other seven lose at most 10. The loss takes the form of a collapse of the output. On the six fragile models, the share of generations that run to the token limit or never state an answer rises by 26 to 78 points, while the share of generations that state a wrong answer does not rise. The collapse is a property of the model. Hermes-3, a different post-training of the same base model as Llama-3.1-8B, loses half as much accuracy as Llama-3.1-8B does. Truncation gains follow the collapse. On the models that hold their accuracy, no truncation sampler improves on plain temperature sampling at 0.7 or 1.0. On the models that collapse, every truncation sampler improves accuracy at 1.3, and at temperatures up to 2.0 truncation delays or prevents the collapse. The gains reported for truncation samplers are therefore real, and they are recoveries from a collapse that only some models suffer and that begins above the temperatures systems use. This work makes three contributions:

- **A controlled comparison of truncation samplers at the temperatures systems use.** Across thirteen open-weight models on GSM8K and MMLU-Pro, no truncation sampler improves on plain temperature sampling at 0.7 or 1.0, and an upper bound states how large a gain the results leave possible.
- **Temperature fragility as a property of the model.** Six of the thirteen models collapse between 1.0 and 1.3 into generations that run to the token limit or never state an answer. The collapse replicates across inference engine and precision. A different post-training of the same base model halves it. A plain

temperature test identifies the fragile models before any sampler is compared.

- **The condition under which truncation samplers help.** On the collapsing models, every truncation sampler improves accuracy at 1.3 and delays the collapse at temperatures up to 2.0. This locates the gains reported for truncation samplers above the temperatures systems use, which reconciles those reports with the absence of gains below.

2 Related work

Truncation samplers. A truncation sampler removes the least probable tokens from the distribution before each token is sampled. Top- k keeps the k most probable tokens (Fan et al., 2018). Top- p , or nucleus sampling, keeps the smallest set of tokens whose probabilities sum to p , so the set grows when the distribution is flat (Holtzman et al., 2020). Min- p keeps every token whose probability is at least a fixed fraction of the largest (Nguyen et al., 2025). Top- $n\sigma$ works on the logits, the model’s scores before they are converted to probabilities, and keeps the tokens within n standard deviations of the largest logit (Tang et al., 2024). Temperature rescales all logits equally, so top- $n\sigma$ keeps the same tokens at every temperature, and so does top- k , because scaling does not change the order of the tokens. Top- p and min- p decide from the probabilities themselves, which temperature changes, so for them the order of scaling and truncation matters. A sampler applied after temperature scaling sees the scaled distribution and keeps a different set of tokens from the same sampler applied before it. Our main configurations apply temperature first, and for top- p and min- p we also run the reverse order, which is the default in llama.cpp.

Temperature-invariant samplers. The parameter of top- p or min- p fixes how much of the distribution is kept, and temperature reshapes the distribution. The same value therefore keeps a different set of tokens at each temperature, and a value that works well at one temperature need not work well at another. Two samplers published in 2026 begin from this observation. p -less sampling (Tan et al., 2026) and Min- k (Ding et al., 2026) each decide which tokens to keep from the shape of the distribution at every step, so that temperature does not change which tokens are kept. Neither paper measures whether that matters at the temperatures systems use. The evaluation of Min- k begins at

temperature 1.0, where top- k , top- p , and min- p score within a point of one another on GSM8K, and the samplers separate only above temperature 2 (Ding et al., 2026).

Evaluations of decoding and temperature. Studies that compare decoding methods find that the best one depends on the model. A survey of decoding methods reports that the best method changes with the task and the model (Shi et al., 2024). The best configuration for open-ended generation varies by model (Garces Arias et al., 2025), and so does the gap between greedy decoding and sampling on reasoning benchmarks (Song et al., 2025). Studies of temperature either stop at 1.0 or use tasks with little accuracy to lose. Renze and Guven (2024) find no effect of temperature between 0.0 and 1.0 on multiple-choice problems. Grover (2026) concludes that instruction tuning is the main determinant of robustness to temperature, on synthetic tasks where the most robust models score 11 to 30% and with no truncation sampler tested. Fastowski et al. (2025) link the entropy of the token distribution to temperature sensitivity on factual questions, and Troshin et al. (2025) attribute high-temperature losses on mathematics to wrong tokens at a few critical positions.

Conditions of measurement. One recent study shapes the design of our comparison. Masoudian et al. (2026) run three instruction-tuned models through five inference frameworks on six benchmarks and find that the framework changes the scores even under greedy decoding, by amounts that depend on the model. A comparison across models and samplers must therefore hold the framework fixed, or the differences it reports may belong to the framework. We run every configuration through one engine, llama.cpp (Gerganov and ggml contributors, 2023). To check whether the collapse we observe comes from the engine, we rerun the most fragile model through a second engine at full precision. Our models also run quantized, with weights stored at reduced precision to fit in less memory. Prasad and Pal (2026) find that quantization does not change how much a model loses to temperature, and our grid at four quantization levels agrees (Appendix C).

3 Experimental design

Table 1 lists the six runs. All of them use the same items, prompts, token budgets, and parser, so a difference in accuracy between two runs comes

from the factor that differs between them, up to the sampling variation that the confidence intervals quantify.

Decoding. Every model in the main grid runs under eight decoding configurations. Two are base-lines, greedy decoding and plain temperature sampling with no truncation. Four apply one truncation sampler after temperature scaling, each at the setting its source paper uses: top- p 0.95 (Holtzman et al., 2020), top- k 40 (Radford et al., 2019), min- p 0.05 (Nguyen et al., 2025), and top- $n\sigma$ 1.0 (Tang et al., 2024). The last two repeat top- p and min- p with truncation applied before temperature scaling, the default order in llama.cpp, so both orders are tested for the two samplers where the order matters (Section 2). Temperatures 0.7 and 1.0 bracket the defaults of Section 1, and 1.3 sits above them, to show what happens when a model is run past its default range.

Tasks. GSM8K is a set of grade-school word problems with a numeric answer (Cobbe et al., 2021). We use the zero-shot chain-of-thought prompt of Meta’s Llama evaluation suite, which asks for reasoning without worked examples (Meta AI, 2024c). MMLU-Pro is ten-option multiple choice across 14 categories (Wang et al., 2024), prompted the same way. Both tasks need a chain of reasoning that ends in one answer that can be checked exactly, and MMLU-Pro’s chains run about twice as long (median 398 against 212 completion tokens at temperature 0.7). Each task uses 50 fixed items, the first 50 of its test split.

Grading. Each generation has a token budget, 512 tokens on GSM8K and 1024 on MMLU-Pro. A generation that is still running when it reaches the budget is cut off there, and we call it *capped*. The parsers follow lm-eval-harness (Gao et al., 2023), the framework behind most published scores on these tasks, so our grading matches the scores readers know and no parsing choice of ours can favour one configuration. The parser records one of three outcomes. A generation is *strict* when it contains the final-answer sentence the prompt asks for (“The final answer is 26”; “The answer is (I)”). It is *flexible* when that sentence is absent and the parser falls back to the last number or the last parenthesized letter. It is *failed* when neither exists, and it is then graded incorrect. Appendix B shows one generation of each kind.

Models. The main grid holds seven instruction-tuned models from five families: Llama-3.1-8B (Grattafiori et al., 2024), Mistral-7B-v0.3 (Jiang

Table 1: The six runs, 232,700 graded generations in total. Every run uses both tasks with the same fixed 50 items each and $K=3$ repetitions of every sampled configuration ($K=1$ for greedy decoding), except the stratified subset, which uses 50 fresh MMLU-Pro items. “8” configurations are the eight of the Decoding paragraph; K is the number of repetitions per item. The main grid includes a BF16 GGUF run of Qwen3-4B beside its four quantized runs. Gen. counts graded generations.

Run	Models	Precision, engine	Configurations	T	Gen.	Where
Main grid	7 (5 families)	Q8, Q6, Q4, Q3; llama.cpp	8	0.7, 1.0, 1.3	185,600	§4.1–4.4
Temperature extension	6 more	Q8; llama.cpp	greedy, plain temp.	0.7, 1.0, 1.3	6,000	§4.1, 4.2
Flagged-model grid	3 of the 6	Q8; llama.cpp	8	0.7, 1.0, 1.3	19,200	§4.3, 4.4
Engine and release checks	Llama-3.2, Llama-3.1, Qwen3-4B	BF16, INT8; transformers. F16, Q8; llama.cpp	greedy, plain temp.	0.7, 1.0, 1.3	6,000	§4.1, App. A
Temperature ladder	Llama-3.1, Qwen2.5, Gemma-3	Q8; llama.cpp	4	1.3, 1.5, 1.7, 2.0	14,400	§4.4
Stratified subset	Llama-3.1, Qwen2.5, Gemma-3	Q8; llama.cpp	greedy, plain temp.	0.7, 1.0, 1.3	1,500	App. C

et al., 2023), Qwen2.5-7B (Yang et al., 2024), Qwen3-1.7B, 4B, and 8B with thinking off (Yang et al., 2025), and Gemma-3-12B (Gemma Team, 2025). They are the widely used open-weight families at sizes that fit the single 16 GB consumer GPU on which every run was made (Appendix A). Each runs at four quantization levels, Q8_0, Q6_K, Q4_K_M, and Q3_K_M, where the digit is the number of bits per weight, all produced by one provider with the same conversion settings (Appendix A), so the levels differ only in precision. Every temperature result in the paper uses Q8_0, the level closest to full precision. Six more models run the temperature test at Q8_0. Two are further Llama-3 releases, Llama-3-8B (Meta AI, 2024a) and Llama-3.2-3B (Meta AI, 2024b), added because Llama-3.1-8B collapsed in the grid and we wanted to know whether the other releases of the family do. Hermes-3-8B (Teknium et al., 2024) is a different post-training of the Llama-3.1-8B base, which separates the base model from its post-training. Qwen3.5-9B (Qwen Team, 2026) and Gemma-4-E4B (Google DeepMind, 2026) are two releases from 2026, added to test whether current models show the same pattern, and OLMo-3-7B (Allen Institute for AI, 2025) is a fully open model. Appendix A gives the model files, their hashes, and the llama.cpp builds.

Fragility flag. The six additional models first run the temperature test alone. A model then receives the eight decoding configurations only if the test shows that it loses accuracy to temperature, because the sampler comparison on models that hold their accuracy is already covered by the main grid. We fixed the criterion before any of the six ran. A model qualifies when, between temperatures 0.7 and 1.3 under plain sampling, its accuracy falls by at least 15 points, its capped rate rises by at least 20 points, or its share of strict parses falls by at least 20 points, on either task. Three models qualified, Hermes-3, Qwen3.5-9B, and OLMo-3-7B,

and each ran the eight configurations at Q8_0.

Engine check. Every result in the grid comes from llama.cpp on quantized weights, so a collapse could come from the engine or from the conversion of the weights. To check, we rerun the temperature test through Hugging Face Transformers (Wolf et al., 2020) on the official weights: Llama-3.2-3B at full 16-bit precision, Llama-3.1-8B in 8-bit weights, the largest that fit the GPU, and Qwen3-4B at full precision as a model that holds its accuracy.

Statistical inference. All intervals in the paper are 95% confidence intervals obtained by bootstrapping over items. We resample the 50 items in each task, keeping all generations of an item together. This captures uncertainty due to the choice of items without treating repeated generations as independent observations. Comparisons between configurations are paired because each configuration is evaluated on the same 50 items. In Section 4.3, we ask whether any truncation sampler improves on plain temperature sampling. To account for testing several samplers, we resample their differences from temperature sampling jointly and compute an upper confidence bound on the largest improvement (Westfall and Young, 1993).

4 Results

4.1 Temperature sensitivity differs by model

Figure 1 gives the accuracy drop from temperature 0.7 to 1.3 under plain sampling for all thirteen models at Q8_0, with the accuracies at every temperature in Appendix C. The models divide into two groups. Six lose between 17 and 38 points on MMLU-Pro. The other seven models lose at most 10 points, and for each of these seven the confidence interval includes a drop of 4 points or less. We call the first group fragile and the second robust. On GSM8K, the shorter task, only Llama-3.1-8B and Llama-3.2-3B lose heavily, 30 and 38 points, and every other model loses 12 points or less.




Figure 1: Accuracy drop under plain temperature sampling from $T=0.7$ to 1.3 at Q8_0, with 95% item-clustered bootstrap intervals; the accuracies at the two temperatures are printed at the right. Positive values are losses. Models are sorted by the MMLU-Pro drop; the six in orange above the dotted line lose at least 15 points on one task, the accuracy criterion of the fragility flag.

The collapse survives a change of engine and precision. Run through Hugging Face Transformers on its official 16-bit weights, Llama-3.2-3B loses 40 points on GSM8K and 36 on MMLU-Pro, against 38 and 37 through llama.cpp at Q8_0, a difference within the confidence intervals on both tasks. Llama-3.1-8B in 8-bit weights collapses as well, and Qwen3-4B at full precision holds its accuracy as it does at Q8_0 (Appendix A). Post-training changes the size of the collapse. Hermes-3-8B shares the Llama-3.1-8B base and differs from it only in post-training, and it loses half as much on MMLU-Pro, 19 points against 38, with the same rise in capped generations.

4.2 The lost accuracy is degenerate output

An accuracy drop can come from two kinds of failure. Either the model completes an answer and the answer is wrong, or the model never reaches an answer. The parser separates the two. Figure 2 divides every MMLU-Pro generation under plain temperature sampling into three outcomes. The first is correct. The second is capped or unparseable: the generation reached its token budget, or it contained no answer the parser could grade. The third is answered wrong: the generation ended before the budget with a parsed answer, and the answer is wrong. In all six fragile models the capped-or-unparseable share rises between temperatures 0.7 and 1.3, by 26 to 78 points, while

the answered-wrong share falls or stays flat. In the seven robust models the capped-or-unparseable share moves by at most 11 points. GSM8K shows the same pattern for Llama-3.1-8B and Llama-3.2-3B (Appendix B), with one difference in grading. Incoherent text usually contains a number, so on GSM8K the parser’s fallback grades it as a wrong answer, while on MMLU-Pro it fails to parse. The capped flag records the failure on both tasks.

Capped generations at 0.7 and at 1.3 fail in different ways. To see what they contain, we sampled 152 generations from four models across the three outcomes and the three temperatures and had Claude Fable 5.1 (Anthropic, 2026) label each one as completed, misformatted, cut off while coherent, or degenerate. The labelling model saw only the task and the raw text, and it did not know which model or temperature had produced the text. Every audited Llama-3.1 generation capped at 1.3 was labelled degenerate. Every one capped at 0.7 was labelled cut off while coherent. The second kind is common at 0.7 in verbose models: OLMo-3 and Qwen3.5 reach the budget on 47 and 29% of their MMLU-Pro generations there.

4.3 No sampler gains where accuracy holds

On the six robust models with a sampler grid, no truncation sampler improves on plain temperature sampling at temperatures 0.7 and 1.0. Table 2 summarizes the comparisons. Each truncation sampler




Figure 2: Outcome of every MMLU-Pro generation under plain temperature sampling at Q8_0, three bars per model, top to bottom $T=0.7$, $T=1.0$, and $T=1.3$. Categories are exclusive, in this precedence: correct; capped or unparseable (so this category holds incorrect generations only); answered wrong. Model order as in Figure 1.

is compared with plain temperature sampling on the same items, for every robust model in the main grid, both tasks, and both temperatures. Almost every difference lies within 3 points of zero, and the largest gain is 3.3 points. Because the question is whether any of these many comparisons shows a gain, the table also gives an upper confidence bound on the largest gain across all of them, which is 9.0 points. The bound is a limit on how much a truncation sampler could gain on these models and tasks at these temperatures.

The first gains appear where accuracy has begun to fall. On the same six models at temperature 1.3, gains reach 9 points, and the comparisons with intervals above zero belong to Mistral-7B, Qwen2.5-7B, and Gemma-3-12B, the three robust models with the largest drops. Llama-3.1-8B, the fragile model in the main grid, shows a gain already at 1.0. Its MMLU-Pro accuracy has fallen 10 points by that temperature, and min- p and top- $n\sigma$ each recover about 11 of them. Reversing the order of temperature scaling and truncation changes accuracy by at most 4 points on the main-grid models at 0.7 and 1.0 (Appendix E).

Table 2: Paired gains of the four truncation rules over plain temperature, by family of contrasts: the number of contrasts, the largest point gain, the approximate 95% upper bound on the largest gain (the 95th percentile of the jointly resampled maximum; not computed for the families at 1.3, where the gains are the finding), and the number of pointwise intervals above zero. The six robust grid models are Mistral-7B, Qwen2.5-7B, Qwen3-1.7B, Qwen3-4B, Qwen3-8B, and Gemma-3-12B. Grid rows pool quantization levels; panel rows are Q8_0.

Models	T	contrasts	largest gain (pp)	95% bound (pp)	intervals > 0
six robust grid models	0.7, 1.0	96	+3.3	9.0	1
six robust grid models, Q8_0 only	0.7, 1.0	96	+4.7	16.0	0
six robust grid models	1.3	48	+9.0	—	13
Llama-3.1-8B	0.7, 1.0	16	+11.2	16.8	6
three flagged panel models	0.7, 1.0	48	+8.0	18.0	2
three flagged panel models	1.3	24	+21.3	32.7	20

4.4 Truncation recovers accuracy on the collapse

On the fragile models, every truncation sampler improves accuracy at temperature 1.3 (Figure 3). The gains are large because the collapse is large. On every model and task the best sampler brings accuracy back to within six points of its level at 0.7, so the accuracy lost between 0.7 and 1.3 under plain sampling was almost entirely recoverable by removing the tail of the distribution. Which sampler recovers the most varies from model to model, and because each sampler was compared with plain temperature sampling only, we do not rank them.

Where each model collapses. The robust models collapse too, at higher temperatures. Figure 4 extends the temperature to 2.0 on Llama-3.1-8B, Qwen2.5-7B, and Gemma-3-12B under plain sampling and three truncation samplers. A model has collapsed once its accuracy falls below half of what it was at 0.7, and the collapse temperature is the first tested temperature at which that happens. Under plain sampling the three models collapse in the same order as in Figure 1, Llama-3.1 first and Gemma-3 last, and each collapses earlier on MMLU-Pro, the longer task, than on GSM8K. Truncation moves the collapse to a higher temperature or removes it from the tested range. Top- p delays it in half of the model and task combinations, min- p delays it further, and top- $n\sigma$ does not collapse on any model or task through 2.0. Min- p and top- $n\sigma$ were validated at temperatures of 1.5 to 3 (Nguyen et al., 2025; Tang et al., 2024), the range in which these collapses occur, and that is where their gains are largest here as well.




Figure 3: Truncation at $T=1.3$ on the four fragile checkpoints with a sampler grid (Llama-3.1-8B at Q8_0 from the main grid; the three flagged panel models). Rules at their standard settings with temperature applied first. Left: paired difference from plain temperature at 1.3. Right: paired difference from plain temperature at 0.7. Pointwise 95% item-clustered bootstrap intervals. Every difference in the left panel is positive, and on the three flagged models 20 of the 24 intervals exclude zero. For the flagged models both plain-temperature baselines are rerun within the sampler-grid run and differ from the extension run’s by up to 6 points.




Figure 4: Accuracy against temperature under four configurations at Q8_0 (points at 0.7 and 1.0 from the main grid, 1.3 to 2.0 from the temperature sweep). The dashed line marks half of the same model and task’s plain-temperature accuracy at 0.7 in the main grid, the collapse criterion; crossings are read from point estimates on the tested temperatures.

5 Discussion and conclusion

Under plain temperature sampling, some checkpoints keep their accuracy above the default tem-

peratures and others lose it by 1.3, and the loss appears as capped or unparseable generations. Larger plain-temperature losses accompany larger truncation gains; the recovery contrasts of Figure 3 carry that claim on their own, and the association across all cells is descriptive. This reconciles the sampler literature with the small gains we measure at default temperatures. Min- p and top- $n\sigma$ were evaluated in a range that overlaps the collapse temperatures we observe, and there they help; Min- k ’s comparison at temperature 1.0 shows the probability-based rules within a point of each other (Ding et al., 2026).

For a practitioner running objective tasks, the two settings should be chosen together. On a model that holds its accuracy at the intended temperature, the gain from a standard truncation rule is small: across 96 contrasts on the six robust grid models at 0.7 and 1.0, the largest estimated gain is 3.3 points and the simultaneous upper bound 9.0 (16.0 at Q8_0 alone). On a model that loses accuracy there, either a lower temperature or a truncation rule recovers most of the loss, and the fragility can be measured without gold answers: the capped rate and the strict-parse rate need no labels, and their change between 0.7 and 1.3 separates the fragile models from the robust ones in every run of this study, including the panel run after the flag was written. We measure accuracy only, so we say nothing about the output diversity a higher temperature buys.

For sampler evaluations, a gain should be reported next to the plain-temperature loss it recovers, since a gain measured beyond one model’s collapse temperature can be absent on a model that collapses later. What moves the collapse temperature remains open. Hermes-3 shows that post-training changes its size on one base, and the grid contains no base-versus-instruct or checkpoint axis, so the cause is not identified here.

Limitations

Each task has 50 fixed items with $K=3$ repetitions; the small item count is what makes a factorial run of 185,600 generations feasible, and every main estimate carries an item-clustered interval. The fixed MMLU-Pro subset is single-category (business) because the test split is sorted by category; the stratified subset preserves Llama’s much larger drop on 50 fresh items, and as a second sample of 50 it does not estimate the benchmark-wide effect

either. Mistral-7B scores 31% on GSM8K and 29% on MMLU-Pro at 0.7, which limits the power to detect a further drop on a model we count as robust.

Temperature 1.3 is a stress point above the defaults, chosen to bound the deployment range from above. The token budgets are fixed, and budget censoring is visible at 0.7 on the verbose models; the audit distinguishes cut-off reasoning from degenerate text on a sample of 152 outputs from four grid models, labelled by a language model with 34 borderline cases; the remaining records, including every output of the three flagged panel models, are unaudited. The survivor comparison controls for item composition with matched greedy decoding; it does not remove the selection.

The sampler evidence is narrower than the temperature evidence. Thirteen models ran the temperature test; ten ran a sampler grid, and the three flagged models were selected by their temperature drop, so the recovery estimates describe a fragility-defined subgroup. “Best rule” is chosen after the fact, which favours the recovery. Rules were tested at one setting each, in one engine, and the bound applies to those settings. The two engine checks replicate the collapse for two Llama checkpoints and one control; Llama-3.1’s check ran in 8-bit weights, its native precision untested, and the intervals on the between-engine differences are about 20 points wide. Repeated generations, temperatures, and rules are not independent observations, so the gain-versus-loss correlation is descriptive.

The scope is open-weight instruction-tuned models of 1.7B to 12B parameters, in English, zero-shot, on objective tasks with one checkable answer, with thinking mode off. A spot-check of Qwen3-8B with thinking on shows a 10-point MMLU-Pro drop [2, 20] at $N=25$ (Appendix C), so the bound does not cover thinking mode. Open-ended generation, the original use of truncation samplers, lies outside the design, as does the diversity a higher temperature buys under multi-sample selection. The study identifies no cause for the differences between checkpoints.

References

- Allen Institute for AI. 2025. OLMo 3 7b instruct model card. <https://huggingface.co/allenai/olmo-3-7B-Instruct>.
- Anthropic. 2026. Claude Fable 5.1. <https://www.anthropic.com/claude/fable>. Model page, read September 2026.

579	Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian,	Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and	631
580	Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias	Yejin Choi. 2020. The curious case of neural text de-	632
581	Plappert, Jerry Tworek, Jacob Hilton, Reiichiro	generation . In <i>International Conference on Learning</i>	633
582	Nakano, Christopher Hesse, and John Schulman.	<i>Representations (ICLR)</i> .	634
583	2021. Training verifiers to solve math word prob-		
584	lems . <i>arXiv preprint arXiv:2110.14168</i> .		
585	Yuanhao Ding, Meimingwei Li, Esteban Garces Arias,	Albert Q. Jiang, Alexandre Sablayrolles, Arthur Men-	635
586	Matthias Aßenmacher, Christian Heumann, and	sch, Chris Bamford, Devendra Singh Chaplot, Diego	636
587	Chongsheng Zhang. 2026. Min-k sampling: De-	de las Casas, Florian Bressand, Gianna Lengyel, Guil-	637
588	coupling truncation from temperature scaling via rel-	laume Lample, Lucile Saulnier, et al. 2023. Mistral	638
589	ative logit dynamics . In <i>Proceedings of the 64th An-</i>	7B . <i>arXiv preprint arXiv:2310.06825</i> .	639
590	<i>annual Meeting of the Association for Computational</i>		
591	<i>Linguistics (Volume 1: Long Papers)</i> , pages 14932–	Shahed Masoudian, Passant Shafaei, Monorama Swain,	640
592	14948.	and Markus Schedl. 2026. What we observe as LLM	641
		behavior can be a side-effect of inference backend .	642
		<i>arXiv preprint arXiv:2608.04714</i> .	643
593	Angela Fan, Mike Lewis, and Yann Dauphin. 2018.	Meta AI. 2024a. Introducing Meta Llama 3: The most	644
594	Hierarchical neural story generation . In <i>Proceedings</i>	capable openly available LLM to date. https://ai.	645
595	<i>of ACL</i> .	meta.com/blog/meta-llama-3/ .	646
596	Alina Fastowski, Bardh Prenkaj, and Gjergji Kasneci.	Meta AI. 2024b. Llama 3.2: Revolutionizing	647
597	2025. From confidence to collapse in LLM factual	edge AI and vision with open, customizable mod-	648
598	robustness . <i>arXiv preprint arXiv:2508.16267</i> .	els. https://ai.meta.com/blog/llama-3-2-	649
		connect-2024-vision-edge-mobile-devices/ .	650
599	Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman,	Meta AI. 2024c. Meta Llama 3 evaluation de-	651
600	Sid Black, et al. 2023. A framework for few-shot	tails. https://github.com/meta-llama/llama3/	652
601	language model evaluation . Zenodo.	blob/main/eval_details.md .	653
602	Esteban Garces Arias, Meimingwei Li, Christian	Minh Nhat Nguyen, Andrew Baker, Clement Neo,	654
603	Heumann, and Matthias Aßenmacher. 2025. Decod-	Allen Roush, Andreas Kirsch, and Ravid Shwartz-	655
604	ing decoded: Understanding hyperparameter effects	Ziv. 2025. Turning up the heat: Min-p sampling for	656
605	in open-ended text generation . In <i>Proceedings of</i>	creative and coherent LLM outputs . In <i>International</i>	657
606	<i>the 31st International Conference on Computational</i>	<i>Conference on Learning Representations (ICLR)</i> .	658
607	<i>Linguistics (COLING)</i> .		
608	Gemma Team. 2025. Gemma 3 technical report . <i>arXiv</i>	Hari Prasad and Ritam Pal. 2026. The joint effect	659
609	<i>preprint arXiv:2503.19786</i> .	of quantization and sampling temperature on LLM	660
610	Georgi Gerganov and ggml contributors. 2023.	safety alignment: A factorial analysis . <i>arXiv preprint</i>	661
611	llama.cpp: LLM inference in C/C++. https://	<i>arXiv:2606.29581</i> .	662
612	github.com/ggml-org/llama.cpp . Software.		
613	Google DeepMind. 2026. Gemma 4 E4B instruction-	Qwen Team. 2026. Qwen3.5-9b model card. https:	663
614	tuned model card. https://huggingface.co/	//huggingface.co/Qwen/Qwen3.5-9B .	664
615	google/gemma-4-E4B-it .		
616	Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri,	Alec Radford, Jeffrey Wu, Rewon Child, David Luan,	665
617	et al. 2024. The Llama 3 herd of models . <i>arXiv</i>	Dario Amodei, and Ilya Sutskever. 2019. Language	666
618	<i>preprint arXiv:2407.21783</i> .	models are unsupervised multitask learners . <i>OpenAI</i>	667
		<i>technical report</i> .	668
619	Kabir Grover. 2026. Reliability under randomness:	Matthew Renze and Erhan Guven. 2024. The effect of	669
620	An empirical analysis of sparse and dense language	sampling temperature on problem solving in large	670
621	model behavior across decoding temperatures . <i>arXiv</i>	language models . In <i>Findings of EMNLP</i> .	671
622	<i>preprint arXiv:2601.00942</i> .		
623	John Hewitt, Christopher D. Manning, and Percy Liang.	Rylan Schaeffer, Joshua Kazdan, and Yegor Denisov-	672
624	2022. Truncation sampling as language model	Blanch. 2025. Min-p, max exaggeration: A criti-	673
625	desmoothing . In <i>Findings of EMNLP</i> .	cal analysis of min-p sampling in language models .	674
		<i>arXiv preprint arXiv:2506.13681</i> .	675
626	Andreas Hochlehnert, Hardik Bhatnagar, Vishaal Udan-	Chufan Shi, Haoran Yang, Deng Cai, Zhisong Zhang,	676
627	darao, Samuel Albanie, Ameya Prabhu, and Matthias	Yifan Wang, Yujiu Yang, and Wai Lam. 2024. A	677
628	Bethge. 2025. A sober look at progress in language	thorough examination of decoding methods in the era	678
629	model reasoning: Pitfalls and paths to reproducibility .	of LLMs . In <i>Proceedings of EMNLP</i> .	679
630	<i>arXiv preprint arXiv:2504.07086</i> .	Yifan Song, Guoyin Wang, Sujian Li, and Bill Yuchen	680
		Lin. 2025. The good, the bad, and the greedy: Eval-	681
		uation of LLMs should not ignore non-determinism .	682
		In <i>Proceedings of NAACL</i> .	683

Runyan Tan, Shuang Wu, and Phillip Howard. 2026. *p*-less sampling: A robust hyperparameter-free approach for LLM decoding. In *International Conference on Learning Representations (ICLR)*.

Chenxia Tang, Jianchun Liu, Hongli Xu, and Liusheng Huang. 2024. *Top- σ : Not all logits are you need*. *arXiv preprint arXiv:2411.07641*.

Ryan Teknium, Jeffrey Quesnelle, and Chen Guang. 2024. *Hermes 3 technical report*. <https://arxiv.org/abs/2408.11857>. *Preprint*, arXiv:2408.11857.

Sergey Troshin, Irina Mohammed, Yan Meng, Christof Monz, Antske Fokkens, and Vlad Niculae. 2025. *Control the temperature: Selective sampling for diverse and high-quality LLM outputs*. In *Conference on Language Modeling (COLM)*.

Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhui Chen. 2024. *MMLU-Pro: A more robust and challenging multi-task language understanding benchmark*. In *Advances in Neural Information Processing Systems*.

Peter H. Westfall and S. Stanley Young. 1993. *Resampling-Based Multiple Testing: Examples and Methods for p-Value Adjustment*. Wiley, New York.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, and 3 others. 2020. *Transformers: State-of-the-art natural language processing*. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45.

An Yang et al. 2024. *Qwen2.5 technical report*. *arXiv preprint arXiv:2412.15115*.

An Yang et al. 2025. *Qwen3 technical report*. *arXiv preprint arXiv:2505.09388*.

A Reproducibility and engine checks

Build and item checks. Llama-3.2-3B and Qwen3-4B rerun on the later llama.cpp release reproduce their pinned-build drops with differences that include zero, and the stratified MMLU-Pro subset preserves Llama’s much larger drop (Llama-3.1 +43.3 [31.3, 55.3]; Qwen2.5 +6.0; Gemma-3 +7.3).

Inference stack. llama.cpp at commit 5aba5364 (server build 9456), native /completion endpoint, CUDA builds per GPU architecture (sm_75 for an RTX 2080 Ti, sm_120 for an RTX 5070 Ti). Server-side BOS insertion is disabled; the client-rendered

chat template owns the special tokens, and each record stores the SHA-256 of the rendered prompt. The 2026 panel, its sampler grids, and the bridge run use release v0.4.0 (commit 5266f24d), which two panel models require, with the same client-side rendering and the server’s template parser off. That release forces its own BOS for Gemma 4, so the Gemma-4 prompt is rendered without one; on both builds the server returns some cap-length degenerate streams with a zero token count, the pinned build with their text (which the client recovers, and which the audit labelled) and v0.4.0 as an HTTP error without text. The runner records those as zero-token capped generations. They number 189 of 185,600 records in the main grid, 114 in the Llama-releases run, 158 in the panel, 196 in its sampler grids, 64 in the engine check, 73 in the bridge, and 2,549 on the ladder; at $T=1.3$ under plain temperature they are 13% of Llama-3.1’s capped generations. Every quantized checkpoint is a Bartowski imatrix GGUF; SHA-256 hashes of the model files are in the repository (DOWNLOADS.md).

Decoding and seeding. Exactly one truncation parameter per request; the chain order is pinned through the server’s samplers array and recorded in every record. At both builds the server instantiates only the samplers named in that array, so its defaults for unlisted samplers (top- k 40, top- p 0.95, min- p 0.05) are inert, and no penalty, DRY, XTC, or logit-bias parameter is sent. Every generation’s seed is derived from (global seed 0, model, quantization, sampler, temperature, task, item, repetition). Batched serving (≤ 8 slots) is seed-logged but not bit-exact. Tasks: GSM8K (openai/gsm8k, test split, first 50 items) and MMLU-Pro (TIGER-Lab/MMLU-Pro, test split, first 50 items). Each generation is one JSONL record with the identity tuple, request parameters, raw output, parse outcome, and token counts; runs are resumable by record id. The pipeline, configurations, analysis scripts, and every record are at <https://anonymous.4open.science/r/decoding-robustness-027C>.

Engine check. The transformers configuration loads meta-llama/Llama-3.2-3B-Instruct (revision 0cb88a4f) in BF16 with transformers 5.16.1 and torch 2.14, samples with do_sample at temperature T with top- k and top- p disabled and no other processor, in batches of eight with left padding; greedy decoding is do_sample=False. Generation stops on the model’s end-of-sequence set or at the task’s budget. The prompt is ren-

dered from the same template as the grid; the template stamps the run date into the Llama system header, the only difference between the runs’ prompt strings, and the prompt token counts agree on every item. Llama-3.1-8B loads the unsloth/Llama-3.1-8B-Instruct mirror of the official weights (revision 4699cc75) in 8-bit through bitsandbytes, in batches of four; Qwen3-4B loads Qwen/Qwen3-4B (revision 1cfa9a72) in BF16 with thinking disabled. Both reproduce every grid prompt hash. Table 3 gives accuracy levels, drops, capped rates, and the paired difference of each configuration’s drop from the pinned Q8_0 llama.cpp cell, including the v0.4.0 bridge rows.

B Grading and degeneration

Surviving answers at 1.3. Among the generations that do complete with the pinned final-answer sentence, accuracy at 1.3 is subject to selection, because the surviving generations can come from a different mix of items. The matched comparison is greedy accuracy on the surviving generations’ items, each survivor counted once. On MMLU-Pro, Llama-3.1’s completed answers score 78.9% against 78.9% for matched greedy, Qwen3.5’s 83.1 against 84.7, and OLMo-3’s 73.3 against 73.3; for Llama-3-8B, Llama-3.2-3B (nine survivors), and Hermes-3 the completed answers score below matched greedy, by 4, 11, and 16 points (Appendix B). Three fragile checkpoints score 4 to 16 points below matched greedy among their surviving answers, in point estimates, and so does robust Mistral-7B (21 points), so a survivor gap at 1.3 is not specific to fragility (Table 5 gives the 0.7 reference for every model).

Figure 5 shows one generation per parse outcome and Figure 6 two degenerate outputs. The strict rule matches the pinned sentence with the answer directly after it. The flexible rule catches two different things: a well-formed answer written outside the pinned sentence, and an incoherent generation that still contains a number, which GSM8K grades as a wrong answer. Mistral omits the pinned sentence on 26 to 32% of its GSM8K generations at every temperature, and the Qwen3 models wrap the answer in markdown emphasis, which the strict rule does not match; both grade correctly through the flexible rule. Because markdown defeats the strict rule, we re-graded every main-grid record with a strict rule that tolerates it: 598 of 185,600 grades change, 556 of them on Qwen3-8B and

```
[strict] ... x = 19.50 / 0.75    x = 26    So
the original price of the book was $26.    The
final answer is $26.
```

```
[flexible] ... 80 \text{ miles} + 150 \text{ miles} = 230 \text{ miles}    ### Final Answer:
The final answer is 230.
```

```
[failed] ... \frac{3,462.20}{35} = 98.92
### Final Answer:    The answer is J.
```

Figure 5: One generation per parse outcome (endings only). Strict: Llama-3.1-8B, GSM8K, $T=1.3$, graded 26, correct. Flexible: Qwen3-8B, GSM8K, $T=0.7$; the emphasis markers defeat the strict rule and the last-number fallback grades 230, correct. Failed: Qwen3-8B, MMLU-Pro, $T=0.7$; the parser expects a parenthesized letter such as (J), none appears, graded incorrect.

```
[opens] To solve this problem, we'll use the
gross profit method of inventory evaluation.
Step 1: Calculate the gross profit earned
during the year. . .
[ends] H sieve Con Palette sediment conduit
France called coaching OMAX may PHYS today
Enhanced vertices XOR/t and registration
ordinances efforts yesterday Michael
```

```
[opens] To determine the correct type of
research methods, we need to understand the
characteristics of each option: A.
Non-probability: This method involves
selecting. . .
[ends] entageeba lever modeling ign measures
inherits cones quantum auditingippines image
later campaigns Taylorart parte diesel El N
resembled Dental professionals
```

Figure 6: Two degenerate outputs (Llama-3.1-8B, Q8_0, plain temperature, $T=1.3$, MMLU-Pro; both reached the 1024-token budget). Each opens coherently, loses coherence mid-generation, and continues as token noise until the budget.

Qwen3-4B MMLU-Pro, and no plain-temperature drop moves by more than 1.0 point. The original parser, fixed before any run, stays primary. Table 4 gives the share of each outcome and the capped rate for every model and temperature, and Figure 7 the GSM8K outcome composition.

Blinded audit. To check that the parser outcomes mean what Section 4.2 takes them to mean, we drew 152 plain-temperature generations at Q8_0 from Llama-3.1-8B, Mistral-7B, Qwen2.5-7B, and Qwen3-8B at $T \in \{0.7, 1.3\}$ on both tasks, stratified by parse outcome with the uncertain outcomes oversampled, and had them labelled by a language-model annotator (Claude) that saw only the task and the raw output: completed, cut off while coherent, degenerate, misformatted, or other.

Table 3: Engine, precision, and release checks: plain-temperature accuracy at $T=0.7$ and 1.3, the drop with its 95% item-clustered interval, the capped rate at 1.3 (%), and the difference of the drop from the reference Q8_0 llama.cpp cell on the same items (positive: a larger drop than the reference). The BF16 GGUF row of Qwen3-4B belongs to the main grid.

Model	Precision, engine	GSM8K					MMLU-Pro				
		0.7	1.3	drop [CI]	cap	vs. ref. [CI]	0.7	1.3	drop [CI]	cap	vs. ref. [CI]
Llama-3.2-3B	BF16, transformers	74.7	34.7	+40.0 [29.3, 49.3]	61	+2.0 [-8.7, 12.0]	40.7	4.7	+36.0 [24.7, 48.0]	92	-0.7 [-9.3, 8.7]
	F16 GGUF, llama.cpp	74.0	42.7	+31.3 [22.7, 40.7]	53	-6.7 [-17.3, 4.0]	38.0	5.3	+32.7 [23.3, 43.3]	91	-4.0 [-10.7, 2.7]
	Q8 GGUF, llama.cpp v0.4.0	77.3	34.7	+42.7 [34.0, 50.7]	63	+4.7 [-3.3, 12.7]	39.3	4.7	+34.7 [24.7, 45.3]	93	-2.0 [-9.3, 5.3]
	Q8 GGUF, llama.cpp (reference)	74.0	36.0	+38.0 [28.7, 46.7]	55	-	40.0	3.3	+36.7 [27.3, 46.7]	92	-
Llama-3.1-8B	INT8, transformers	82.7	39.3	+43.3 [33.3, 52.0]	54	+13.3 [2.7, 23.3]	45.3	2.7	+42.7 [32.7, 53.3]	95	+4.7 [-3.3, 12.7]
	Q8 GGUF, llama.cpp (reference)	77.3	47.3	+30.0 [21.3, 39.3]	51	-	48.0	10.0	+38.0 [28.0, 48.0]	86	-
Qwen3-4B	BF16, transformers	94.7	93.3	+1.3 [-1.3, 4.0]	0	-1.3 [-5.3, 2.0]	60.7	61.3	-0.7 [-6.0, 5.3]	13	+0.0 [-6.7, 6.7]
	BF16 GGUF, llama.cpp	95.3	96.7	-1.3 [-4.0, 1.3]	0	-4.0 [-9.3, 1.3]	58.0	60.0	-2.0 [-7.3, 3.3]	15	-1.3 [-8.7, 6.0]
	Q8 GGUF, llama.cpp v0.4.0	96.7	94.0	+2.7 [-0.7, 7.3]	0	+0.0 [-4.0, 4.7]	65.3	64.0	+1.3 [-4.7, 7.3]	11	+2.0 [-6.0, 10.0]
	Q8 GGUF, llama.cpp (reference)	96.0	93.3	+2.7 [-1.3, 6.7]	1	-	64.0	64.7	-0.7 [-6.0, 4.7]	12	-




Figure 7: Outcome of every GSM8K generation under plain temperature sampling at Q8_0, as in Figure 2. On GSM8K a degenerate generation usually contains a number, so part of the degeneration is graded through the flexible path; the capped flag records it.

Of the 51 capped outputs, all 8 of Llama’s at 1.3 are degenerate and all 8 at 0.7 are coherent reasoning cut short; the robust models’ 21 capped outputs at 1.3 split 8 degenerate, 9 cut off, 3 completed, and 1 misformatted, and their 14 at 0.7 split 13 cut off and 1 completed. Of the 26 parse failures, 15 contain a clear answer in an unexpected form, 6 are degenerate, 4 are refusals, and 1 is cut off. Of the 75 strict and flexible parses, 72 are completed answers. The annotator marked 34 of the 152 labels as borderline. Sample, labels, and key are in the repository (results/outcome_audit).

Survivor comparison. Table 5 gives, for every model, the share of plain-temperature genera-

tions with a strict parse at 0.7 and 1.3, the accuracy among them, and greedy accuracy on the same surviving items.

C Temperature results in full

Table 6 gives plain-temperature accuracy at all three temperatures for the thirteen models, the drop with its interval, and the change in capped and strict-parse rates. Table 7 recomputes the drop at each quantization level of the main grid: Llama-3.1 collapses at every level and no robust model exceeds 12 points at any level. Quantization itself costs little greedy accuracy: from Q8_0 to Q3_K_M, greedy accuracy pooled over both tasks changes by at most 5 points on the models of 7B parameters or more and falls by 7 and 9 points on Qwen3-4B and Qwen3-1.7B. Table 8 repeats the temperature test on a stratified MMLU-Pro subset with 3 to 4 items from each of the 14 categories; Llama’s drop stays far above the other two, and Gemma-3’s stratified drop, +7.3 [2.0, 12.7], has an interval that excludes zero. A spot-check of Qwen3-8B with thinking on (Q8_0, $N=25$, $K=2$, 450 generations) reproduces the GSM8K result (94% at 0.7 against 96% at 1.3) and not the MMLU-Pro one (56 against 46%, a 10-point drop with interval [2, 20]); with thinking on, the longer traces make Qwen3-8B less robust on the longer task.

D Temperature ladder

Table 9 gives the plain-temperature drop at each ladder temperature against the model’s own accuracy at 0.7, with intervals. The ladder is a separate run from the main grid, and its 1.3 step differs from Table 6 by amounts within the intervals (Llama-3.1 on MMLU-Pro: +44.7 against +38.0).

Table 4: Parse outcomes under plain temperature at Q8_0 (% of generations, rounded): strict, flexible, failed, and capped (capped overlaps the other three).

Model	T	GSM8K				MMLU-Pro			
		strict	flex.	failed	capped	strict	flex.	failed	capped
Llama-3.1-8B	0.7	89	11	0	9	79	1	19	11
	1.0	87	13	0	7	63	5	32	22
	1.3	45	55	0	51	13	6	81	86
Llama-3.2-3B	0.7	99	1	0	1	83	3	15	10
	1.0	89	11	0	2	65	5	30	22
	1.3	41	59	0	55	6	9	85	92
Qwen3.5-9B	0.7	80	20	0	21	71	3	26	29
	1.0	79	21	0	21	67	5	27	33
	1.3	71	29	0	27	39	9	51	43
Hermes-3-8B	0.7	68	32	0	0	93	2	5	0
	1.0	68	32	0	0	91	1	9	0
	1.3	59	39	2	9	57	2	41	38
OLMo-3-7B	0.7	84	16	0	14	56	11	33	47
	1.0	84	16	0	16	62	10	28	38
	1.3	78	17	5	22	30	0	70	68
Llama-3-8B	0.7	98	2	0	0	97	1	2	0
	1.0	98	2	0	0	97	1	2	0
	1.3	94	6	0	5	65	4	31	29
Mistral-7B	0.7	72	28	0	5	82	4	14	1
	1.0	73	27	0	6	81	2	17	1
	1.3	71	29	0	7	74	6	20	11
Qwen2.5-7B	0.7	97	3	0	0	98	1	1	2
	1.0	93	7	0	0	96	1	3	3
	1.3	93	7	0	1	88	3	9	9
Gemma-4-E4B	0.7	45	55	0	25	74	8	18	26
	1.0	54	46	0	28	67	8	25	34
	1.3	49	51	0	29	68	10	22	32
Gemma-3-12B	0.7	97	3	0	3	91	1	9	9
	1.0	98	2	0	2	89	3	9	11
	1.3	95	5	0	5	79	3	17	20
Qwen3-8B	0.7	45	55	0	1	82	1	17	12
	1.0	45	55	0	2	83	5	13	13
	1.3	47	53	0	2	78	5	17	18
Qwen3-1.7B	0.7	15	85	0	1	87	2	11	13
	1.0	22	78	0	1	85	1	14	14
	1.3	21	79	0	4	83	2	15	19
Qwen3-4B	0.7	96	4	0	1	89	0	11	10
	1.0	93	7	0	1	87	1	12	12
	1.3	96	4	0	1	88	1	11	12

E Sampler results in full

Loss and gain across cells. Across the 60 model, task, and temperature combinations with a sampler

Table 5: Accuracy among strict-parsed generations under plain temperature at Q8_0, with the share of generations that parse strictly, the number n of such generations at $T=1.3$ (of 150), and greedy accuracy on their items with each surviving generation counted once. Point estimates without intervals.

Model	Task	$T=0.7$			$T=1.3$			
		strict %	acc strict	greedy, same items	strict %	n	acc strict	greedy, same items
Llama-3.1-8B	GSM8K	89	85.1	81.3	45	67	98.5	95.5
	MMLU-Pro	79	60.5	56.3	13	19	78.9	78.9
Llama-3.2-3B	GSM8K	99	74.3	74.3	41	62	85.5	90.3
	MMLU-Pro	83	47.6	48.4	6	9	55.6	66.7
Qwen3.5-9B	GSM8K	80	97.5	92.5	71	106	99.1	93.4
	MMLU-Pro	71	86.8	81.1	39	59	83.1	84.7
Hermes-3-8B	GSM8K	68	86.3	83.3	59	89	83.1	82.0
	MMLU-Pro	93	50.0	54.3	57	86	47.7	64.0
OLMo-3-7B	GSM8K	84	98.4	97.6	78	117	97.4	95.7
	MMLU-Pro	56	69.0	61.9	30	45	73.3	73.3
Llama-3-8B	GSM8K	98	76.9	76.9	94	141	75.2	77.3
	MMLU-Pro	97	45.2	39.7	65	98	41.8	45.9
Mistral-7B	GSM8K	72	33.3	26.9	71	106	32.1	26.4
	MMLU-Pro	82	35.0	39.8	74	111	25.2	45.9
Qwen2.5-7B	GSM8K	97	94.5	93.8	93	139	91.4	93.5
	MMLU-Pro	98	66.7	63.3	88	132	65.2	68.2
Gemma-4-E4B	GSM8K	45	98.5	94.0	49	74	93.2	90.5
	MMLU-Pro	74	83.8	81.1	68	102	82.4	83.3
Gemma-3-12B	GSM8K	97	93.8	90.4	95	143	93.0	93.0
	MMLU-Pro	91	69.9	62.5	79	119	75.6	67.2
Qwen3-8B	GSM8K	45	97.0	98.5	47	70	100.0	100.0
	MMLU-Pro	82	74.0	70.7	78	117	73.5	74.4
Qwen3-1.7B	GSM8K	15	91.3	91.3	21	31	71.0	80.6
	MMLU-Pro	87	55.0	55.0	83	124	57.3	54.8
Qwen3-4B	GSM8K	96	95.8	97.9	96	144	93.8	97.9
	MMLU-Pro	89	72.2	72.2	88	132	73.5	73.5

grid (seven grid models at Q8_0 and three flagged models), larger losses of plain temperature against greedy decoding accompany larger gains of truncation over plain temperature (Pearson $r=0.89$ over the 60 best-rule points; 0.78 over the 240 rule-level points). The association is descriptive, and the two quantities share the plain-temperature baseline (Appendix E).

Order of temperature and truncation. The two ablations are outside these families. Applying top- p or min- p before temperature scaling changes accuracy by at most 4 points on the seven main-grid models at $T \leq 1.0$. The largest gain over plain temperature anywhere at $T \leq 1.0$ is the temperature-last top- p on Hermes-3 at 0.7 on MMLU-Pro, +12.0 [4.7, 19.3], one contrast among 48 ablation contrasts on the flagged models that otherwise stay within 7 points. On Llama-3.1 at 1.3, truncating first keeps 22 more points under top- p on MMLU-Pro: at that temperature the rule cuts a sharper distribution and removes more of the tail (Appendix E).

Tables 10 and 11 give the accuracy of every configuration at every temperature for the ten models with a sampler grid (main-grid models pooled over

quantization levels; flagged models at Q8_0). Tables 12 and 13 give every paired contrast against plain temperature, including the two temperature-last ablations; an interval endpoint within 10^{-9} of zero counts as touching zero. A generalized estimating equation over the same records (identity link, item clusters, exchangeable working correlation) gives the same point estimates and intervals that differ by at most 0.6 points (median 0.1). Figure 8 plots the gain of each rule against the loss of plain temperature relative to greedy decoding, 240 points over 60 combinations. The two quantities share the plain-temperature baseline, so the association is descriptive.

Table 6: Plain-temperature accuracy at Q8_0 for all thirteen models, the drop from $T=0.7$ to 1.3 with its 95% item-clustered interval (bold: lower limit at or above 15, a stricter criterion than the fragility flag), and the change in capped and strict-parse rates. Models sorted as in Figure 1.

Model	Task	accuracy (%) at T			drop [95% CI]	change 0.7→1.3 (pp)	
		0.7	1.0	1.3		capped	strict
Llama-3.1-8B	GSM8K	77.3	80.0	47.3	+ 30.0 [21.3, 39.3]	+43	-45
	MMLU-Pro	48.0	38.0	10.0	+ 38.0 [28.0, 48.0]	+75	-67
Llama-3.2-3B	GSM8K	74.0	69.3	36.0	+ 38.0 [28.7, 46.7]	+55	-57
	MMLU-Pro	40.0	33.3	3.3	+ 36.7 [27.3, 46.7]	+82	-77
Qwen3.5-9B	GSM8K	81.3	82.0	72.0	+9.3 [3.3, 16.0]	+6	-9
	MMLU-Pro	61.3	59.3	32.7	+ 28.7 [18.0, 40.0]	+13	-31
Hermes-3-8B	GSM8K	77.3	74.7	65.3	+12.0 [4.0, 20.0]	+9	-9
	MMLU-Pro	46.7	36.0	27.3	+19.3 [10.0, 29.3]	+38	-36
OLMo-3-7B	GSM8K	86.7	82.7	79.3	+7.3 [1.3, 14.7]	+8	-6
	MMLU-Pro	39.3	40.7	22.0	+17.3 [10.0, 25.3]	+21	-26
Llama-3-8B	GSM8K	75.3	74.7	71.3	+4.0 [-3.3, 11.3]	+5	-4
	MMLU-Pro	44.0	42.7	27.3	+16.7 [8.7, 24.7]	+29	-32
Mistral-7B	GSM8K	31.3	28.7	32.0	-0.7 [-9.3, 8.0]	+1	-1
	MMLU-Pro	29.3	34.7	20.0	+9.3 [0.7, 18.0]	+9	-8
Qwen2.5-7B	GSM8K	92.7	87.3	88.7	+4.0 [0.7, 8.0]	+1	-4
	MMLU-Pro	65.3	66.0	58.0	+7.3 [-1.3, 16.0]	+7	-10
Gemma-4-E4B	GSM8K	76.0	73.3	73.3	+2.7 [-1.3, 6.7]	+3	+5
	MMLU-Pro	63.3	55.3	58.7	+4.7 [-0.0, 10.0]	+6	-6
Gemma-3-12B	GSM8K	92.0	91.3	88.7	+3.3 [-1.3, 8.7]	+2	-2
	MMLU-Pro	63.3	64.7	60.0	+3.3 [-4.0, 10.7]	+11	-11
Qwen3-8B	GSM8K	92.7	94.0	94.7	-2.0 [-5.3, 1.3]	+1	+2
	MMLU-Pro	60.7	60.0	58.0	+2.7 [-4.7, 9.3]	+6	-4
Qwen3-1.7B	GSM8K	79.3	75.3	74.0	+5.3 [-2.0, 13.3]	+3	+5
	MMLU-Pro	48.0	47.3	47.3	+0.7 [-7.3, 8.7]	+6	-5
Qwen3-4B	GSM8K	96.0	96.7	93.3	+2.7 [-1.3, 6.7]	+0	+0
	MMLU-Pro	64.0	60.7	64.7	-0.7 [-6.0, 4.7]	+2	-1

Table 7: Plain-temperature accuracy drop $T=0.7 \rightarrow 1.3$ (pp, point estimates) at each quantization level of the main grid (150 generations per cell).

Model	Task	Q8_0	Q6_K	Q4_K_M	Q3_K_M
Llama-3.1-8B	GSM8K	+30.0	+32.0	+36.0	+38.0
	MMLU-Pro	+38.0	+46.7	+38.0	+40.7
Mistral-7B	GSM8K	-0.7	-0.7	+3.3	+12.0
	MMLU-Pro	+9.3	+5.3	+10.7	+6.7
Qwen2.5-7B	GSM8K	+4.0	+2.7	+2.7	+6.7
	MMLU-Pro	+7.3	+6.0	+4.0	+2.0
Gemma-3-12B	GSM8K	+3.3	+2.7	-0.0	+2.0
	MMLU-Pro	+3.3	+5.3	+2.7	+3.3
Qwen3-8B	GSM8K	-2.0	+1.3	+2.0	+0.7
	MMLU-Pro	+2.7	-0.7	+0.0	+2.0
Qwen3-4B	GSM8K	+2.7	+0.7	+6.0	+2.7
	MMLU-Pro	-0.7	+0.0	+2.0	+4.0
Qwen3-1.7B	GSM8K	+5.3	-6.0	-0.7	+0.0
	MMLU-Pro	+0.7	+6.0	-8.0	+4.0

Table 8: Stratified MMLU-Pro subset against the fixed business subset: plain-temperature accuracy at Q8_0 and the drop with its interval.

Model	Subset	acc. 0.7	acc. 1.3	drop [95% CI]
Llama-3.1-8B	business (head)	48.0	10.0	+38.0 [28.0, 48.0]
	stratified	44.7	1.3	+43.3 [31.3, 55.3]
Qwen2.5-7B	business (head)	65.3	58.0	+7.3 [-1.3, 16.0]
	stratified	54.7	48.7	+6.0 [-2.0, 14.7]
Gemma-3-12B	business (head)	63.3	60.0	+3.3 [-4.0, 10.7]
	stratified	64.0	56.7	+7.3 [2.0, 12.7]

Table 9: Plain-temperature accuracy drop at each ladder temperature against the model’s own $T=0.7$ accuracy in the main grid (pp, 95% item-clustered intervals; Q8_0).

Task	T	Llama-3.1-8B	Qwen2.5-7B	Gemma-3-12B
GSM8K	1.3	+32.7 [24.7, 41.3]	+4.0 [-0.7, 9.3]	+2.7 [-2.0, 8.0]
	1.5	+68.0 [58.7, 76.7]	+24.0 [15.3, 32.7]	+14.7 [8.7, 22.0]
	1.7	+76.0 [66.0, 85.3]	+70.0 [61.3, 77.3]	+42.7 [34.7, 50.7]
	2.0	+77.3 [67.3, 86.7]	+90.7 [82.0, 97.3]	+84.0 [76.0, 90.7]
MMLU-Pro	1.3	+44.7 [34.6, 55.3]	+8.0 [0.0, 17.3]	+3.3 [-2.7, 10.0]
	1.5	+46.0 [35.3, 56.7]	+40.0 [28.6, 51.3]	+39.3 [29.3, 50.0]
	1.7	+48.0 [36.7, 59.3]	+61.3 [49.3, 73.3]	+57.3 [46.0, 69.3]
	2.0	+48.0 [36.7, 59.3]	+65.3 [53.3, 77.3]	+63.3 [51.3, 75.3]




Figure 8: Gain of each truncation rule over plain temperature against the loss of plain temperature relative to greedy decoding: 240 points, four rules in each of the 60 model, task, and temperature combinations with a sampler grid (Q8_0). Both axes subtract the same plain-temperature accuracy. The diagonal marks full recovery to greedy accuracy.

Table 10: GSM8K accuracy (%) by configuration and temperature (main-grid models pooled over quantization levels; the three flagged models at Q8_0).

Model	Configuration	0.7	1.0	1.3	Model	Configuration	0.7	1.0	1.3
Llama-3.1-8B	greedy	78.0	–	–	Qwen3-4B	greedy	95.2	–	–
	temperature	78.2	74.8	44.2		temperature	95.2	94.0	93.1
	top- p	81.5	78.0	63.8		top- p	95.5	95.5	93.5
	top- k	79.8	75.2	58.7		top- k	94.0	94.5	94.1
	min- p	81.2	78.0	72.0		min- p	95.1	94.7	92.9
	top- $n\sigma$	78.8	80.7	80.8		top- $n\sigma$	94.8	93.9	93.9
	top- p (temp. last)	80.8	78.7	73.0		top- p (temp. last)	94.8	95.2	94.3
	min- p (temp. last)	78.7	81.3	76.2		min- p (temp. last)	94.5	93.9	93.9
Mistral-7B	greedy	31.0	–	–	Qwen3-1.7B	greedy	75.5	–	–
	temperature	32.7	30.0	29.2		temperature	73.3	73.7	73.7
	top- p	28.2	31.0	34.2		top- p	73.8	75.0	72.7
	top- k	31.8	31.7	29.8		top- k	72.8	72.7	73.3
	min- p	32.3	28.3	28.8		min- p	72.2	71.7	74.0
	top- $n\sigma$	29.5	30.7	30.2		top- $n\sigma$	74.3	73.5	72.3
	top- p (temp. last)	29.5	31.2	29.5		top- p (temp. last)	73.0	72.0	71.3
	min- p (temp. last)	29.5	32.0	30.5		min- p (temp. last)	74.7	73.3	74.5
Qwen2.5-7B	greedy	93.5	–	–	Hermes-3-8B	greedy	82.0	–	–
	temperature	91.5	91.7	87.5		temperature	71.3	77.3	60.0
	top- p	92.2	91.2	90.3		top- p	74.7	80.0	71.3
	top- k	90.2	88.5	88.7		top- k	74.0	72.7	67.3
	min- p	90.8	90.5	91.0		min- p	75.3	76.0	71.3
	top- $n\sigma$	92.7	91.2	90.7		top- $n\sigma$	70.7	74.7	77.3
	top- p (temp. last)	90.8	91.3	90.3		top- p (temp. last)	72.0	73.3	77.3
	min- p (temp. last)	92.0	90.7	91.3		min- p (temp. last)	75.3	77.3	74.7
Gemma-3-12B	greedy	92.0	–	–	Qwen3.5-9B	greedy	78.0	–	–
	temperature	91.0	90.2	89.0		temperature	81.3	80.0	74.0
	top- p	91.5	90.3	90.7		top- p	80.7	80.7	80.0
	top- k	91.8	91.7	90.7		top- k	81.3	80.7	77.3
	min- p	91.5	91.3	92.3		min- p	82.7	82.7	83.3
	top- $n\sigma$	91.3	91.5	91.0		top- $n\sigma$	78.7	84.7	81.3
	top- p (temp. last)	90.2	91.8	91.5		top- p (temp. last)	82.7	80.7	82.0
	min- p (temp. last)	91.2	90.8	90.3		min- p (temp. last)	82.0	78.0	79.3
Qwen3-8B	greedy	94.5	–	–	OLMo-3-7B	greedy	90.0	–	–
	temperature	95.0	94.2	94.5		temperature	88.7	81.3	76.0
	top- p	94.0	94.3	93.7		top- p	84.7	88.7	82.0
	top- k	94.8	94.8	93.5		top- k	84.7	85.3	78.7
	min- p	94.8	95.3	94.7		min- p	84.7	83.3	78.0
	top- $n\sigma$	96.3	94.2	94.3		top- $n\sigma$	86.7	84.0	86.0
	top- p (temp. last)	93.8	94.0	94.0		top- p (temp. last)	84.0	85.3	81.3
	min- p (temp. last)	96.2	94.7	95.3		min- p (temp. last)	83.3	84.7	86.0

Table 11: MMLU-Pro accuracy (%) by configuration and temperature (main-grid models pooled over quantization levels; the three flagged models at Q8_0).

Model	Configuration	0.7	1.0	1.3	Model	Configuration	0.7	1.0	1.3
Llama-3.1-8B	greedy	54.0	—	—	Qwen3-4B	greedy	61.6	—	—
	temperature	47.8	39.7	7.0		temperature	61.9	62.5	61.2
	top- p	51.3	46.0	16.0		top- p	62.8	61.3	62.8
	top- k	51.2	43.3	26.8		top- k	63.5	61.1	61.2
	min- p	51.2	50.8	39.5		min- p	60.9	62.7	62.9
	top- $n\sigma$	51.0	50.5	48.8		top- $n\sigma$	62.4	62.4	61.5
	top- p (temp. last)	50.3	46.7	37.5		top- p (temp. last)	62.9	62.7	61.7
	min- p (temp. last)	51.2	50.2	43.5		min- p (temp. last)	61.3	63.6	61.5
Mistral-7B	greedy	35.5	—	—	Qwen3-1.7B	greedy	45.5	—	—
	temperature	30.5	30.2	22.5		temperature	46.0	45.7	45.3
	top- p	30.2	29.5	25.3		top- p	47.2	47.0	47.2
	top- k	32.5	32.3	28.3		top- k	47.5	46.5	45.0
	min- p	33.8	30.7	31.0		min- p	47.0	47.7	46.2
	top- $n\sigma$	32.3	31.5	31.5		top- $n\sigma$	46.8	47.2	46.2
	top- p (temp. last)	29.2	30.2	32.5		top- p (temp. last)	47.7	46.8	43.5
	min- p (temp. last)	33.2	33.3	28.7		min- p (temp. last)	45.7	46.2	49.0
Qwen2.5-7B	greedy	64.0	—	—	Hermes-3-8B	greedy	52.0	—	—
	temperature	62.5	61.3	57.7		temperature	42.7	42.7	26.0
	top- p	65.5	62.8	62.7		top- p	41.3	47.3	38.7
	top- k	62.5	61.7	57.0		top- k	42.7	42.0	34.7
	min- p	61.3	63.5	59.7		min- p	44.0	42.0	46.0
	top- $n\sigma$	62.0	63.5	61.0		top- $n\sigma$	44.7	50.7	38.7
	top- p (temp. last)	63.7	62.7	61.3		top- p (temp. last)	54.7	46.7	43.3
	min- p (temp. last)	64.2	63.0	62.3		min- p (temp. last)	44.0	44.7	44.0
Gemma-3-12B	greedy	61.0	—	—	Qwen3.5-9B	greedy	64.0	—	—
	temperature	63.7	66.7	60.0		temperature	60.0	58.7	38.0
	top- p	65.7	66.7	63.8		top- p	62.0	60.7	53.3
	top- k	64.3	62.8	63.7		top- k	59.3	56.7	53.3
	min- p	65.3	67.0	65.5		min- p	61.3	56.0	55.3
	top- $n\sigma$	65.2	65.8	65.8		top- $n\sigma$	58.0	61.3	59.3
	top- p (temp. last)	65.8	65.0	65.2		top- p (temp. last)	56.7	60.7	54.0
	min- p (temp. last)	65.3	63.7	63.8		min- p (temp. last)	60.0	57.3	57.3
Qwen3-8B	greedy	60.5	—	—	OLMo-3-7B	greedy	46.0	—	—
	temperature	60.5	59.0	59.5		temperature	42.0	41.3	26.7
	top- p	60.5	61.3	60.5		top- p	43.3	39.3	36.7
	top- k	60.2	61.3	57.0		top- k	40.0	40.7	37.3
	min- p	61.0	59.2	59.7		min- p	43.3	40.7	42.7
	top- $n\sigma$	60.0	58.8	58.7		top- $n\sigma$	41.3	42.7	40.0
	top- p (temp. last)	60.7	59.5	60.5		top- p (temp. last)	43.3	44.0	44.7
	min- p (temp. last)	61.3	60.7	59.3		min- p (temp. last)	42.0	44.7	40.0

Table 12: GSM8K: paired accuracy difference of each configuration against plain temperature on the same items (pp, 95% item-clustered bootstrap intervals; main-grid models pooled over quantization, flagged models at Q8_0). Bold: interval above zero.

Model	T	top- p	top- k	min- p	top- $n\sigma$	top- p (temp. last)	min- p (temp. last)
Llama-3.1-8B	0.7	+ 3.3 [0.3, 6.8]	+1.7 [-1.0, 4.5]	+3.0 [-0.8, 7.0]	+0.7 [-2.7, 4.0]	+2.7 [-0.5, 6.3]	+0.5 [-2.3, 3.3]
	1.0	+3.2 [-0.3, 6.8]	+0.3 [-3.5, 4.5]	+3.2 [-0.8, 7.3]	+ 5.8 [2.2, 9.5]	+ 3.8 [0.2, 8.0]	+ 6.5 [2.8, 10.5]
	1.3	+ 19.7 [15.0, 24.3]	+ 14.5 [10.0, 18.8]	+ 27.8 [22.7, 33.2]	+ 36.7 [30.3, 43.2]	+ 28.8 [23.7, 34.2]	+ 32.0 [26.0, 38.3]
Mistral-7B	0.7	-4.5 [-7.7, -1.5]	-0.8 [-5.8, 3.7]	-0.3 [-3.8, 3.3]	-3.2 [-7.7, 1.0]	-3.2 [-8.0, 1.5]	-3.2 [-7.2, 0.7]
	1.0	+1.0 [-2.5, 4.7]	+1.7 [-1.5, 4.7]	-1.7 [-4.8, 1.7]	+0.7 [-2.8, 4.2]	+1.2 [-3.0, 5.5]	+2.0 [-1.8, 6.0]
	1.3	+ 5.0 [0.7, 9.2]	+0.7 [-4.3, 5.5]	-0.3 [-4.0, 3.5]	+1.0 [-2.8, 5.0]	+0.3 [-4.8, 5.0]	+1.3 [-2.3, 5.0]
Qwen2.5-7B	0.7	+0.7 [-1.3, 2.8]	-1.3 [-3.8, 0.7]	-0.7 [-2.0, 0.7]	+1.2 [-0.7, 3.2]	-0.7 [-2.8, 1.3]	+0.5 [-1.0, 2.2]
	1.0	-0.5 [-3.0, 1.7]	-3.2 [-6.0, -0.5]	-1.2 [-3.5, 1.2]	-0.5 [-2.8, 1.5]	-0.3 [-3.3, 2.5]	-1.0 [-3.7, 1.5]
	1.3	+ 2.8 [0.7, 5.2]	+1.2 [-0.7, 3.3]	+ 3.5 [0.8, 6.5]	+ 3.2 [1.2, 5.5]	+ 2.8 [0.5, 5.5]	+ 3.8 [1.5, 6.7]
Gemma-3-12B	0.7	+0.5 [-1.3, 2.5]	+0.8 [-1.0, 3.0]	+0.5 [-1.3, 2.5]	+0.3 [-1.5, 2.5]	-0.8 [-4.0, 2.2]	+0.2 [-2.0, 2.5]
	1.0	+0.2 [-3.0, 3.0]	+1.5 [-0.5, 3.7]	+1.2 [-1.2, 3.7]	+1.3 [-1.0, 3.7]	+1.7 [-1.0, 4.7]	+0.7 [-1.5, 2.8]
	1.3	+1.7 [-0.2, 3.8]	+1.7 [-0.7, 4.2]	+ 3.3 [0.5, 6.7]	+2.0 [-0.0, 4.5]	+ 2.5 [0.3, 4.8]	+1.3 [-0.8, 3.5]
Qwen3-8B	0.7	-1.0 [-3.3, 0.8]	-0.2 [-1.7, 1.0]	-0.2 [-1.7, 1.2]	+1.3 [-0.8, 3.5]	-1.2 [-3.5, 0.8]	+1.2 [-0.8, 3.5]
	1.0	+0.2 [-1.7, 2.0]	+0.7 [-2.0, 3.3]	+1.2 [-1.0, 3.3]	-0.0 [-1.8, 1.7]	-0.2 [-2.0, 2.0]	+0.5 [-1.8, 2.8]
	1.3	-0.8 [-2.5, 0.7]	-1.0 [-2.7, 0.5]	+0.2 [-1.8, 2.3]	-0.2 [-1.7, 1.7]	-0.5 [-2.2, 1.0]	+0.8 [-0.3, 2.2]
Qwen3-4B	0.7	+0.3 [-0.9, 1.7]	-1.2 [-2.9, 0.8]	-0.1 [-1.3, 1.3]	-0.4 [-1.7, 1.1]	-0.4 [-2.0, 1.6]	-0.7 [-2.5, 1.6]
	1.0	+1.5 [0.0, 3.1]	+0.5 [-0.5, 1.7]	+0.7 [-1.1, 2.7]	-0.1 [-1.5, 1.1]	+1.2 [0.0, 2.4]	-0.1 [-1.6, 1.3]
	1.3	+0.4 [-1.2, 2.0]	+1.1 [-0.7, 2.8]	-0.1 [-1.9, 1.6]	+0.8 [-1.1, 2.7]	+1.2 [-0.5, 3.1]	+0.8 [-0.8, 2.5]
Qwen3-1.7B	0.7	+0.5 [-2.8, 3.8]	-0.5 [-4.2, 2.8]	-1.2 [-4.7, 2.2]	+1.0 [-2.2, 4.3]	-0.3 [-3.0, 2.2]	+1.3 [-1.3, 4.3]
	1.0	+1.3 [-2.0, 4.5]	-1.0 [-4.0, 2.0]	-2.0 [-4.3, -0.0]	-0.2 [-2.7, 2.3]	-1.7 [-4.0, 0.5]	-0.3 [-3.2, 2.7]
	1.3	-1.0 [-4.5, 2.3]	-0.3 [-3.5, 2.7]	+0.3 [-2.8, 3.5]	-1.3 [-4.2, 1.2]	-2.3 [-6.0, 1.2]	+0.8 [-2.2, 3.7]
Hermes-3-8B	0.7	+3.3 [-3.3, 10.0]	+2.7 [-3.3, 8.7]	+4.0 [-2.0, 10.0]	-0.7 [-6.0, 4.7]	+0.7 [-6.0, 7.3]	+4.0 [-2.7, 10.0]
	1.0	+2.7 [-3.3, 8.0]	-4.7 [-10.7, 1.3]	-1.3 [-6.7, 4.0]	-2.7 [-9.3, 3.3]	-4.0 [-10.0, 2.0]	+0.0 [-7.3, 7.3]
	1.3	+ 11.3 [2.7, 19.3]	+7.3 [-0.7, 14.7]	+ 11.3 [3.3, 19.3]	+ 17.3 [10.7, 24.0]	+ 17.3 [8.7, 25.3]	+ 14.7 [7.3, 22.0]
Qwen3.5-9B	0.7	-0.7 [-4.7, 3.3]	-0.0 [-5.3, 5.3]	+1.3 [-3.3, 5.3]	-2.7 [-8.0, 2.0]	+1.3 [-3.3, 5.3]	+0.7 [-2.7, 4.7]
	1.0	+0.7 [-5.3, 6.7]	+0.7 [-6.0, 7.3]	+2.7 [-1.3, 7.3]	+4.7 [-1.3, 11.3]	+0.7 [-4.7, 6.7]	-2.0 [-6.0, 1.3]
	1.3	+ 6.0 [1.3, 11.3]	+3.3 [-4.0, 10.7]	+ 9.3 [0.7, 18.0]	+ 7.3 [1.3, 13.3]	+ 8.0 [2.0, 14.7]	+5.3 [-1.3, 12.0]
OLMo-3-7B	0.7	-4.0 [-8.7, 0.7]	-4.0 [-9.3, 0.0]	-4.0 [-9.3, 0.7]	-2.0 [-5.3, 1.3]	-4.7 [-10.7, 0.7]	-5.3 [-10.7, -1.3]
	1.0	+ 7.3 [3.3, 12.0]	+4.0 [0.0, 8.7]	+2.0 [-2.7, 6.7]	+2.7 [-3.3, 8.7]	+4.0 [0.0, 8.7]	+3.3 [-0.7, 8.0]
	1.3	+ 6.0 [1.3, 10.7]	+2.7 [-3.3, 9.3]	+2.0 [-4.0, 9.3]	+ 10.0 [3.3, 17.3]	+5.3 [-0.7, 11.3]	+ 10.0 [3.3, 17.3]

Table 13: MMLU-Pro: paired accuracy difference of each configuration against plain temperature on the same items, as in Table 12.

Model	T	top- p	top- k	min- p	top- $n\sigma$	top- p (temp. last)	min- p (temp. last)
Llama-3.1-8B	0.7	+ 3.5 [0.3, 7.0]	+3.3 [-0.5, 7.5]	+3.3 [-0.7, 8.0]	+3.2 [-0.2, 6.7]	+2.5 [-0.7, 6.2]	+ 3.3 [0.3, 6.7]
	1.0	+ 6.3 [1.8, 11.3]	+3.7 [-1.3, 8.7]	+ 11.2 [6.7, 16.0]	+ 10.8 [6.0, 16.2]	+ 7.0 [3.0, 11.2]	+ 10.5 [6.3, 14.7]
	1.3	+ 9.0 [5.2, 13.3]	+ 19.8 [13.8, 26.7]	+ 32.5 [24.7, 41.0]	+ 41.8 [32.3, 52.2]	+ 30.5 [23.2, 38.7]	+ 36.5 [28.0, 45.5]
Mistral-7B	0.7	-0.3 [-4.5, 4.0]	+2.0 [-2.3, 6.5]	+3.3 [-0.5, 7.0]	+1.8 [-2.2, 6.0]	-1.3 [-4.8, 2.2]	+2.7 [-1.7, 7.3]
	1.0	-0.7 [-4.5, 3.0]	+2.2 [-1.2, 5.5]	+0.5 [-2.8, 3.8]	+1.3 [-3.2, 5.8]	-0.0 [-4.0, 4.5]	+3.2 [-0.8, 7.3]
	1.3	+2.8 [-1.0, 7.0]	+ 5.8 [2.7, 9.3]	+ 8.5 [4.5, 12.5]	+ 9.0 [4.8, 13.5]	+ 10.0 [5.3, 15.3]	+ 6.2 [2.0, 10.3]
Qwen2.5-7B	0.7	+ 3.0 [0.5, 5.8]	+0.0 [-2.7, 2.7]	-1.2 [-4.5, 2.2]	-0.5 [-3.7, 2.5]	+1.2 [-1.2, 3.7]	+1.7 [-1.2, 4.7]
	1.0	+1.5 [-1.8, 5.0]	+0.3 [-2.0, 2.7]	+2.2 [-0.8, 5.3]	+2.2 [-0.8, 5.3]	+1.3 [-1.0, 3.8]	+1.7 [-1.2, 4.8]
	1.3	+ 5.0 [1.7, 8.3]	-0.7 [-3.8, 2.3]	+2.0 [-1.7, 6.5]	+3.3 [-0.2, 7.0]	+ 3.7 [0.2, 7.2]	+ 4.7 [1.0, 8.8]
Gemma-3-12B	0.7	+2.0 [-0.3, 4.5]	+0.7 [-2.0, 3.5]	+1.7 [-0.8, 4.3]	+1.5 [-1.2, 4.3]	+2.2 [-1.0, 5.7]	+1.7 [-2.0, 5.7]
	1.0	-0.0 [-2.8, 2.7]	-3.8 [-7.2, -1.0]	+0.3 [-2.2, 2.7]	-0.8 [-3.0, 1.3]	-1.7 [-4.7, 1.3]	-3.0 [-6.2, -0.2]
	1.3	+ 3.8 [0.2, 8.3]	+ 3.7 [0.3, 7.5]	+ 5.5 [2.7, 8.8]	+ 5.8 [2.3, 9.8]	+ 5.2 [2.3, 8.3]	+ 3.8 [0.5, 7.5]
Qwen3-8B	0.7	-0.0 [-2.8, 3.0]	-0.3 [-2.7, 2.0]	+0.5 [-2.3, 3.7]	-0.5 [-2.8, 1.8]	+0.2 [-2.7, 2.8]	+0.8 [-2.2, 3.8]
	1.0	+2.3 [-0.5, 5.2]	+2.3 [-0.7, 5.3]	+0.2 [-3.2, 3.3]	-0.2 [-3.2, 2.7]	+0.5 [-2.3, 3.2]	+1.7 [-1.2, 4.5]
	1.3	+1.0 [-1.8, 3.5]	-2.5 [-5.7, 0.7]	+0.2 [-3.0, 3.5]	-0.8 [-4.2, 2.3]	+1.0 [-1.8, 3.8]	-0.2 [-3.5, 3.0]
Qwen3-4B	0.7	+0.9 [-1.5, 3.1]	+1.6 [-0.5, 3.7]	-0.9 [-3.5, 1.6]	+0.5 [-1.5, 2.5]	+1.1 [-1.1, 3.2]	-0.5 [-3.1, 2.0]
	1.0	-1.2 [-3.2, 0.9]	-1.5 [-3.5, 0.4]	+0.1 [-2.7, 3.2]	-0.1 [-2.7, 2.4]	+0.1 [-2.7, 3.2]	+1.1 [-1.5, 3.7]
	1.3	+1.6 [-1.2, 4.3]	+0.0 [-2.7, 2.5]	+1.7 [-0.7, 4.1]	+0.3 [-2.3, 2.8]	+0.5 [-2.3, 3.2]	+0.3 [-2.1, 2.8]
Qwen3-1.7B	0.7	+1.2 [-2.0, 4.2]	+1.5 [-1.7, 4.7]	+1.0 [-2.5, 4.7]	+0.8 [-2.5, 4.2]	+1.7 [-1.3, 4.8]	-0.3 [-3.7, 2.8]
	1.0	+1.3 [-2.7, 5.0]	+0.8 [-2.2, 4.2]	+2.0 [-1.3, 5.3]	+1.5 [-1.8, 4.8]	+1.2 [-2.3, 4.8]	+0.5 [-3.2, 4.2]
	1.3	+1.8 [-2.8, 7.0]	-0.3 [-4.3, 3.5]	+0.8 [-3.8, 5.5]	+0.8 [-3.5, 5.2]	-1.8 [-6.5, 2.8]	+3.7 [-0.7, 8.7]
Hermes-3-8B	0.7	-1.3 [-8.7, 6.0]	+0.0 [-8.7, 8.0]	+1.3 [-6.0, 8.0]	+2.0 [-4.7, 8.7]	+ 12.0 [4.7, 19.3]	+1.3 [-6.7, 9.3]
	1.0	+4.7 [-2.7, 12.0]	-0.7 [-8.0, 6.7]	-0.7 [-8.7, 7.3]	+ 8.0 [1.3, 14.7]	+4.0 [-1.3, 9.3]	+2.0 [-6.7, 10.7]
	1.3	+ 12.7 [5.3, 20.0]	+ 8.7 [2.0, 16.0]	+ 20.0 [12.0, 29.3]	+ 12.7 [4.0, 22.0]	+ 17.3 [8.7, 26.7]	+ 18.0 [10.0, 26.7]
Qwen3.5-9B	0.7	+2.0 [-1.3, 5.3]	-0.7 [-5.3, 4.0]	+1.3 [0.0, 3.3]	-2.0 [-6.7, 2.0]	-3.3 [-8.0, 0.0]	-0.0 [-4.7, 4.0]
	1.0	+2.0 [-1.3, 5.3]	-2.0 [-6.7, 1.3]	-2.7 [-7.3, 2.0]	+2.7 [-2.0, 7.3]	+2.0 [-4.0, 8.0]	-1.3 [-6.0, 3.3]
	1.3	+ 15.3 [9.3, 22.7]	+ 15.3 [8.0, 24.0]	+ 17.3 [10.0, 26.0]	+ 21.3 [13.3, 30.0]	+ 16.0 [9.3, 23.3]	+ 19.3 [12.0, 28.0]
OLMo-3-7B	0.7	+1.3 [-4.0, 6.7]	-2.0 [-8.7, 4.0]	+1.3 [-4.7, 8.0]	-0.7 [-5.3, 3.3]	+1.3 [-4.0, 6.7]	-0.0 [-6.0, 5.3]
	1.0	-2.0 [-8.7, 4.7]	-0.7 [-8.0, 6.7]	-0.7 [-7.3, 6.0]	+1.3 [-6.0, 8.7]	+2.7 [-4.0, 9.3]	+3.3 [-4.0, 10.7]
	1.3	+ 10.0 [2.7, 17.3]	+ 10.7 [4.0, 17.3]	+ 16.0 [8.0, 24.0]	+ 13.3 [6.0, 22.0]	+ 18.0 [10.7, 26.0]	+ 13.3 [6.6, 20.7]